



Table of contents

1. Overview
2. Contact
3. Disclaimer
4. Citing the use of Retrospective
5. Installation notes
6. Running the software
7. Basic operation



1. Overview

Retrospective is a software utility and graphical user interface to reconstruct self-gated cardiac or respiratory CINE MRI acquired with an MR Solutions preclinical MRI system.

Reconstruction is based on compressed sensing, or on a trained deep learning network, after retrospective sorting of the data into cardiac and/or respiratory time frames, using a navigator signal for cardiac and respiratory synchronization. The left ventricle can be segmented automatically in the reconstructed short-axis CINE images.

2. Contact

Retrospective and this manual is written by:

Gustav Strijkers
Amsterdam UMC
Biomedical Engineering and Physics
Amsterdam, the Netherlands
g.j.striekers@amsterdamumc.nl

3. Disclaimer

The software has been extensively tested using mouse data acquired with an MR Solutions preclinical 7.0T/24cm system.

However, this does not warrant the functions contained in the app will meet your requirements or that the operation of the app will be uninterrupted or error-free.

In case of questions or issues, please contact Gustav Strijkers.

4. Citing the use of Retrospective

Retrospective is free of use. However, proper referencing the use of Retrospective in scientific publications is highly appreciated.

Citations to Retrospective can be done using the following reference ([link](#)):

Daal, M.R.R., Strijkers, G.J., Calcagno, C., Garipov, R.R., Wüst, R.C.I., Hautemann, D., Coolen, B.F. *Quantification of Mouse Heart Left Ventricular Function, Myocardial Strain, and Hemodynamic Forces by Cardiovascular Magnetic Resonance Imaging*. J. Vis. Exp. (171), e62595, doi:10.3791/62595 (2021).



5. Installation notes

Software download

The Matlab app and a Windows standalone version (using the free Matlab runtime engine) can be downloaded from GitHub:

<https://github.com/Moby1971?tab=repositories>

Installation of the Windows standalone version

`Retrospective10_app_installer.exe`

Will install the Matlab runtime engine and the Retrospective app.

Bart toolbox download

The advanced reconstruction options in Retrospective require the Bart reconstruction toolbox, which can be downloaded from:

<https://mrirecon.codeberg.page/>

Retrospective is tested with Bart 1.0.00, whose source is also in the folder functions/bart.



Bart toolbox installation in OSX with MacPorts

- (1) Install the Xcode command line tools

```
$ xcode-select --install
```

- (2) Install MacPorts

From <http://www.macports.org/>. It is recommended to install a newer version of gcc from MacPorts, which step 3 does.

- (3) Install packages

```
$ sudo port install fftw-3-single
$ sudo port install gcc14
$ sudo port install libpng
$ sudo port install openblas
$ sudo port install flock
$ sudo port install gmake
$ sudo port install coreutils
$ sudo port install libblocksruntime
```

- (4) Download and unpack Bart

```
$ wget https://github.com/mrirecon/bart/archive/v1.0.00.tar.gz
$ tar xvfz v1.0.00.tar.gz
$ cd bart-1.0.00
```

- (5) Build and install

```
$ export PATH=/opt/local/libexec/gnubin:$PATH
$ CC=gcc-mp-14 gmake
$ sudo gmake install
```

The compiler name must match the gcc version installed in step 3: gcc14 gives gcc-mp-14, gcc15 gives gcc-mp-15.

- (6) Check that it works

```
$ bart version
```

Bart toolbox installation in OSX with HomeBrew

- (1) Install the Xcode command line tools

```
$ xcode-select --install
```

- (2) Install Homebrew

```
$ /bin/bash -c "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
```

- (3) Install the packages Bart needs

```
$ brew install gcc
$ brew install libpng
$ brew install fftw
$ brew install openblas
$ brew install make
$ brew install llvm libomp
```

The Homebrew package that provides GNU make is called make. It installs the program under the name gmake.



(4) Download and unpack Bart

```
$ wget https://github.com/mrirecon/bart/archive/v1.0.00.tar.gz
$ tar xvfz v1.0.00.tar.gz
$ cd bart-1.0.00
```

(5) Build and install

```
$ CC=gcc-15 MACPORTS=0 gmake
$ sudo gmake install
```

The compiler name must match the version Homebrew installed.

Check it with `ls /opt/homebrew/bin/gcc-*`

(6) Check that it works

```
$ bart version
```

Bart toolbox installation in Windows 11

On Windows, Bart runs inside the Windows Subsystem for Linux (WSL).

(1) Install WSL.

Start a Windows PowerShell as administrator and run:

```
wsl --install
```

This enables the feature, installs WSL 2 and installs Ubuntu in one step. A restart is needed. The first time you start the Linux prompt you will be asked to create a user account.

On older systems, or if `wsl --install` is not recognized, use the manual route instead:

```
Enable-WindowsOptionalFeature -Online -FeatureName Microsoft-Windows-Subsystem-
Linux
```

Then restart, list the available distributions with `wsl -l -o`, install one with `wsl --install -d <distribution name>`, and upgrade to WSL 2. See <https://pureinfotech.com/install-wsl-windows-11/> for details.

(2) Update the Linux distribution

Everything from here on is typed at the Linux prompt, not in PowerShell.

```
$ sudo apt-get update
$ sudo apt-get dist-upgrade
```

(3) Install the packages Bart needs (gcc 12 or newer)

```
$ sudo apt-get install make gcc libfftw3-dev liblapacke-dev libpng-dev
libopenblas-dev gfortran
```

(4) Download and unpack Bart

```
$ wget https://github.com/mrirecon/bart/archive/v1.0.00.tar.gz
$ tar xvfz v1.0.00.tar.gz
$ cd bart-1.0.00
```

(5) Build and install

```
$ make
$ sudo make install
```

If you want GPU support, do not run this yet — follow the next section instead, which builds Bart with CUDA. GPU support is highly recommended.

(6) Check that it works

At the Linux prompt:

```
$ bart version
```

And from Windows, which is how Retrospective calls it:

```
wsl bart version
```

Both should print a version number.

Building Bart with GPU support on Windows (highly recommended)

The reconstruction is considerably faster on a GPU.

(1) Install the NVIDIA driver in Windows

Follow <https://docs.nvidia.com/cuda/wsl-user-guide/index.html>. The Windows driver is what provides GPU access inside WSL — do not install a Linux display driver inside WSL.

(2) Install the CUDA toolkit inside WSL

Get the current installer from <https://developer.nvidia.com/cuda-downloads>, choosing Linux → x86_64 → WSL-Ubuntu.

For CUDA 12.8 this is:

```
$ wget https://developer.download.nvidia.com/compute/cuda/12.8.0/
  local_installers/cuda_12.8.0_570.86.10_linux.run
$ sudo sh cuda_12.8.0_570.86.10_linux.run
```

(3) Find the compute capability of your card

Bart has to be compiled for the GPU generation it will run on. A binary built for the wrong generation compiles without complaint and then fails at the first calculation.

Look your card up at <https://developer.nvidia.com/cuda-gpus> and note its compute capability:

GPU generation	Example cards	Value to use
Pascal	GTX 1060, 1080, Titan X	sm_61
Turing	RTX 2060 – 2080	sm_75
Ampere	RTX 3060 – 3090, A100	sm_80 / sm_86
Ada Lovelace	RTX 4060 – 4090	sm_89
Blackwell	RTX 5070 – 5090	sm_120

This is the same kind of problem as the PyTorch installation described later in this section: which build is correct depends on the card in the machine, and no single value suits every computer.



(4) Set the environment variables

Set these in the same shell you will build in. Adjust `CUDA_BASE` if you installed a different CUDA version, and `GPUARCH_FLAGS` to the value from step 3.

```
$ export CUDA=1
$ export CUDA_BASE=/usr/local/cuda-12.8
$ export CUDA_LIB=lib64
$ export GPUARCH_FLAGS="-arch=sm_61"
$ export PATH=$CUDA_BASE/bin:$PATH
$ export LD_LIBRARY_PATH=$CUDA_BASE/$CUDA_LIB:$LD_LIBRARY_PATH
```

The first four are the variables Bart's Makefile reads. `PATH` makes `nvcc` reachable, and `LD_LIBRARY_PATH` lets the finished program find the CUDA libraries when it runs.

(5) Build

```
$ make -j$(nproc) CUDA=1 CUDA_BASE=$CUDA_BASE CUDA_LIB=$CUDA_LIB
  GPUARCH_FLAGS="$GPUARCH_FLAGS"
$ sudo make install
```

If the build fails at the linking stage with an error about `libstdc++` or `libgcc`, point the linker at the installed gcc explicitly and build again.

Check which version you have with `ls /usr/lib/gcc/x86_64-linux-gnu/` and use that number:

```
$ export LDFLAGS="-L/usr/lib/gcc/x86_64-linux-gnu/13 $LDFLAGS"
```

(6) Test the GPU build

```
$ make utest_gpu
```

How Retrospective finds Bart

In the normal case there is nothing to configure: `sudo make install` puts Bart in `/usr/local/bin`, which is where the app looks.

MacOS

Retrospective reads the environment variable `TOOLBOX_PATH`. When it is not set, it falls back to `/usr/local/bin` and then to `/usr/bin`.

If Bart is installed somewhere else, set `TOOLBOX_PATH` to the folder that contains the bart program. Setting it in your shell profile is not enough — Matlab started from the Dock does not read your shell startup files, so it will not see it. Either start Matlab from a terminal, or set it inside Matlab before starting Retrospective:

```
>> setenv('TOOLBOX_PATH','/usr/local/bin')
```

Windows

Retrospective runs Bart inside WSL, as `wsl bart`. What matters is that bart is on the `PATH` inside your Linux distribution; `TOOLBOX_PATH` is not used on Windows. `sudo make install` takes care of this.

Check from a Windows command prompt with:

```
wsl bart version
```

When Retrospective starts it reports the Bart version it found in the message box. If Bart was not found, the message box says so and the app continues with its own Matlab reconstruction.



Installing Python for deep learning reconstruction and LV segmentation

Two features use Python: the deep learning (DRIM) reconstruction and the LV segmentation. Both are optional; the rest of the app runs without them.

You do not have to install the Python packages yourself. At every start Retrospective checks its Python environment; when there is none, it offers to create it and install what it needs into it. This takes a few minutes and happens only once. Progress is written to the message box, and a full log is kept in `retroPythonSetup.log` in the temporary folder (C:\tmp on Windows).

The environment is created in your home directory:

```
macOS      ~/venv_retro
Windows    $HOME/venv_retro  (inside WSL, not on the Windows drive)
```

The following packages are installed, with fixed versions so that every machine ends up with the same thing:

```
lightning    2.6.5
numpy        2.3.5
scipy        1.17.1
PyYAML       6.0.3
onnxruntime  1.28.0
```

PyTorch is deliberately not in that list, and is handled separately, because which build is correct depends on the machine rather than on a version number. See "Graphics card support" below.

What you need to provide yourself !

On most computers, nothing.

Retrospective looks for a Python it can use, and if it does not find one it installs a Python of its own. Three routes are tried in turn: Conda, if the machine already has it; otherwise a self-contained Python that is downloaded and checked against its published checksum before anything in it is run; otherwise Homebrew, on macOS. None of them needs administrator rights, and none of them disturbs a Python you already use — the one the program installs for itself goes in a folder of its own:

```
macOS      ~/.retro-python
Windows    $HOME/.retro-python  (inside WSL)
```

and nothing else on the machine is touched. You are asked before this happens, and the dialog says where the interpreter will go.

Which Python versions can be used ?

Python 3.11, 3.12 or 3.13.

Below them numpy, scipy and onnxruntime publish nothing that fits the pinned versions, and above them the program has not been tried yet. A Python outside that range is not an error and nothing goes wrong — the program simply does not build on it and uses one of its own instead. It is worth knowing on Windows, where a stock Ubuntu in WSL comes with Python 3.10 and will therefore be passed over. When the program installs a Python for itself it installs 3.12.

PyTorch

If PyTorch is already installed and working on your computer, Retrospective uses it as it is. This is the better case, because a PyTorch that already runs on your hardware is by definition the right build for it. If no PyTorch can be found at all, the program installs one itself.



MacOS

Installing Python together with PyTorch — through Miniconda or Homebrew — gives the best result, because PyTorch then comes from a source that knows your hardware. It is not required. It does not need to be the Python that Matlab finds by default; the program locates a suitable interpreter by itself, and installs one if there is none.

Windows

Python runs inside WSL. See the Bart installation section above for setting up WSL. If WSL is not installed at all, the message box says so: installing it is the one step in the chain the program cannot take for you, because `wsl --install` needs administrator rights and a restart. The rest of the app is unaffected and works normally; the deep learning reconstruction and the LV segmentation are simply switched off.

Inside WSL there is nothing else to prepare. On Ubuntu the module that creates virtual environments is packaged separately, and if it is missing the program works around it — first with `virtualenv`, and failing that by building on a Python of its own. If you would rather give it the short route:

```
$ sudo apt install python3-venv
```

but this is not required, and the program only suggests it in the message box if every other route has failed as well.

Graphics card support

The deep learning features are considerably faster on a graphics card, and on a computer without one the reconstruction may be too slow to be practical.

You do not have to choose the right packages yourself. Every time it starts, Retrospective checks what graphics hardware is present and whether the installed PyTorch and `onnxruntime` can actually use it. If they cannot, it offers to replace them with versions that can, and installs them for you. This is checked at every start, not only when the environment is first built, so a graphics card or driver installed later is picked up the next time you start the program — you do not need to reinstall anything by hand.

What this means in practice:

Computer	Reconstruction uses	Segmentation uses
Apple silicon Mac	the built-in GPU (Metal)	the built-in GPU (CoreML)
NVIDIA graphics card	the card (CUDA)	the card (CUDA)
No graphics card	the processor	the processor

On a computer with an NVIDIA card, the CUDA driver has to be installed and working first — on Windows this means the driver in Windows, with GPU access enabled for WSL, as described in the Bart section above. Retrospective can install the right Python packages, but it cannot install your graphics driver.

If the replacement is offered and you decline, everything still works on the processor, and you will be asked again the next time you start the program.

Checking what was actually used. Retrospective reports the device at the start of every reconstruction, in the message box and in the log file `retroAIdrim.log`, as a line like:

```
GPU available: True (mps), used: True
```

If this says no GPU is available on a computer that has one, the graphics driver is not reachable — on Windows, check that `wsl nvidia-smi` lists your card.



6. Running the software

Running in Matlab 2025a

The Retrospective software can be started from its root directory from the command line.

```
>> retrospective10
```

Retrospective starts in about 12 seconds, or about 20 with the AI module on, which checks its Python first. MATLAB needs about 5.5 GB of memory for a 2D cine of 524 MB and about 5 GB for a 3D cine of 272 MB.

Additional licenses may be required.

```
MATLAB  
Image Processing Toolbox  
Parallel Computing Toolbox  
Signal Processing Toolbox  
Wavelet Toolbox
```

Running the Windows standalone

The Windows standalone version can be run from the start menu or the desktop icon.

Settings and log files

Retrospective keeps its settings in a file, `retroSettings.json`, in a folder of its own for every version of the app:

```
macOS      ~/Library/Application Support/Retrospective/<version>  
Windows    %APPDATA%\Retrospective\<version>
```

The file is created with the standard values when a version first starts, and can be edited by hand; changes take effect at the next start. It holds, among others, the navigator and reconstruction presets, the folders where the file browsers open, the time limits for installing Python, the number of cores NLMF may use (`maxNumCores`), and whether the app checks GitHub for a newer version. A file that cannot be read is set aside as `retroSettings-unreadable-<time>.json` and replaced by the standard values.

Three entries apply to the LV segmentation (step 8): `lvUseGpu`, use the graphics card when it can be used (true); `lvMyoDensity`, the myocardial density for the mass (1.05 g/mL); `lvSmoothFrameSelection`, median filter the blood pool volume curve when choosing the ED and ES frames (true; the volumes are always measured unfiltered).

Every session writes a log file, `retrospective-log-<date and time>.txt`, to the same folder. The log keeps every message, including the ones the message box does not show at the chosen message level. The 50 most recent log files are kept.



7. Basic operation



The Retrospective app operates from a single window with 5 main panels.

Panel 1: Dashboard

This panel contains the task buttons that control the reconstruction process. The dashboard tasks need to be completed from top to bottom. A green light next to the task indicates that a task has been completed. When the light is yellow the task is busy. Red indicates not completed yet.

Panel 2: Navigator / Heart / Respiration / Settings

This panel shows the raw navigator signal, the navigator frequency spectrum, and some advanced navigator analysis parameters.

Panel 3: Navigator analysis

This panel shows the navigator signals filtered for cardiac and respiratory motion, as well as the heart and respiratory rates as function of time.

Panel 4: Sequence / Geometry / Reco / Radial

This panel lists the relevant acquisition and reconstruction parameters.

Panel 5: Reco / Radial

In this panel the reconstructed CINE movie is shown. The Radial sub-panel shows the k-space trajectory of radial and UTE acquisitions.



Step 1: Loading data

Press  Load MRD to load the MRD raw data file *.

The data should include navigator data for cardiac and respiratory synchronization.


For data acquired with multiple receivers (phased-array RF coil) select 1 of the MRD files.

Relevant acquisition parameters will be shown in the Sequence tab of panel 4, including the k-space trajectory: linear, zigzag, user defined, P2ROUD (from the .txt file beside the MRD file), 2D radial, 3D UTE (from a lut*.txt file there) or 3D ZTE.


Step 2: The navigator signal

There are **3** types of navigators, which can be acquired independently and in combination, and in both 2D and 3D imaging.

- **Slice FID navigator** (left): This navigator is the free-induction-decay (green) generated from the rewind of the slice selection gradient.
- **External FID navigator** (middle): The free-induction-decay (green) produced by the rewind of a slice selection gradient in a separate non-spatially encoded slice.
- **External ECHO navigator** (right): The gradient-echo (green) produced in a separate slice with a readout gradient in a single dimension. The resulting navigator is the Fourier transform of the echo and therefore a 1-dimensional image (readout profile) of the object.



Depending on which type of navigator is available, one can choose slice FID, external FID, or external ECHO, separately for heart and respiration, in the Navigator tab. The raw navigator signals and their spectra are shown in the *Heart* and *Respiration* tabs; the filtered signals are in panel 3 (Step 4).



For multi receiver coil data, individual navigators can be inspected with the coil spinner. The app picks the clearest coil or coil combination for heart and respiration.

The respiration peaks in the navigator should be positive. The app turns the navigator the right way up; the Flip switch overrules it when needed.

If cardiac and respiratory oscillations are not clearly visible, change the **primary** point slightly and/or increase **# points** (typically 1 to 5) in the Navigator tab to see whether the quality improves.

When **auto reco** ☒ is enabled, Retrospective watches the MRD folder and reconstructs every new MRD file automatically, for seamless integration with the MR Solutions preclinical imaging software (see Auto reco below).

The *Settings* tab contains 3 presets for navigator filtering of mouse, rat, and zebrafish data. These adapt the values **min/max heart** and **respiration rates**, which determine the search windows for the heart and respiration rate.

The parameter **freq Nx discard** determines the number of samples that are discarded between the end of the navigator and the beginning of the imaging echo. It is active when checked. If an FID artifact is visible in the reconstructed images/movie, increase this number. This parameter is usually only relevant for FLASH sequence version < 710.


Images can be reconstructed during in- or expiration.

Default is reconstruction during expiration, in the intervals between breath-takes. For reconstruction of images during inspiration, tick the *toggle in/expiration* check box.

Discarding k-lines during respiration can be overruled by ticking the *ignore respiration* checkbox.

Error and warning sounds can be switched on or off with the *sounds on/off* checkbox.

Choose black/white or classic colors for the user interface with the radio buttons.





The message level (1 to 3) sets how much the message box shows: 1 shows the essentials, higher levels more detail. The log file keeps every message, whatever the level (section 6).


Under Physio filtering, filter order is the order of the Butterworth filters, and min/max heart rate and min/max resp rate bound the search for the heart and respiration rates; the presets set them. Resp shift (samples) shifts the respiration window, to correct a delay of the respiration navigator. With filter per slice ticked, cardiac and respiratory motion are filtered for each slice separately.

Auto reco watches the folder mrdImportPath of the settings file, or on the scanner computer the MRD folder of the MR Solutions software; without either, it stops. Files already in the folder are left alone, and the newest numbered exam folder is followed. Each new file is loaded, filtered, sorted, reconstructed and exported as MRD (on the scanner computer with SUR and Dicom files); the sounds are switched off, and the window changes color while auto reco runs.

Step 3: Inspecting the navigator power spectrum

The bottom parts of the *Heart* and *Respiration* tabs show the power frequency spectrum of the raw navigator signal.

Typically, the spectrum looks as the figure above, with peaks (and higher harmonics) corresponding to the respiratory and cardiac frequencies. The app automatically determines the most likely heart and respiration frequencies.



The green shaded area indicates the position and width of the Butterworth filters, which are used to filter the raw navigator and extract the periodic cardiac and respiratory motion as function of time.

In case the automatic respiration and/or cardiac frequency detection fails, the **center** and **width** filter values can be adjusted in the heart/respiration filters panel, shown below.

Heart / respiration filters		
	center	width
heart (bpm)	514	100
resp. (bpm)	77	20
window (%)	30	reset

Center and **width** determine the center-frequency and width of the Butterworth filters for navigator filtering.

Window (%) determines the percentage of the data that is discarded during respiratory motion.


Press to reset the values.



Step 4: Filter the navigator

Press  **Filter** to filter the raw navigator signal and extract cardiac and respiratory motion as function of time.

The result may look like this:



The sliders can be used to change the **display range** to inspect the full data set.

Use  **fast**  **slow** buttons to automatically advance the navigator and inspect the data.

The top left figure shows in blue the raw navigator signal and in red the signal filtered for cardiac motion. The gray vertical lines indicate the peaks in the oscillations, which can be used to check the alignment between the red and blue signals.



Bottom left is the heart rate as function of time during the full acquisition. The blue line is the heartbeat to heartbeat variation, whereas the red line is a smoothed trend-line.

The top right figure shows in blue the raw navigator signal and in red the signal filtered for respiratory motion. The gray vertical lines indicate the peaks in the oscillations, which can be used to check the alignment between the red and blue signals. The shaded red areas indicates the part of the data which is discarded during respiratory motion.

Bottom right is the respiration rate as function of time during the full acquisition. The blue line is the respiration to respiration variation, the red line is a smoothed trend-line.

Heart and respiration rates should be smoothly varying over time. If this is not the case, adjust the Butterworth filter settings.

Including / excluding data




In case of irregular respiration or heartbeat part of the data can be included (green) and/or excluded (red) by checking the **Include** or **Exclude** boxes or both. Including or excluding data is not possible for multi-slice 2D data.

After this, press  **Filter** again to filter the data with the new settings.



Step 5: Sort the k-space

Choose the desired sorting method, number of frames, and number of dynamics. Reconstruction presets are in the Reco tab of panel 4.



Sorting methods

(1) Cardiac

Cardiac phase-binning. Each heartbeat is divided into an equal number of cardiac CINE frames. The average time between cardiac frames is determined from the average heart-rate and the number of cardiac frames. This option is preferred for most applications, particularly for high frame-rate reconstructions for analysis of diastolic function.

(2) Cardiac abs.

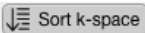
Cardiac absolute-binning. A fixed frame time is predetermined from the average heart-rate and the number of cardiac CINE frames. This option could be more accurate for reconstruction of systolic function in the presence of heart-rate variations, but leads to temporal blurring in the diastolic phase.

(3) Respiratory

Respiratory phase-binning. Each respiration cycle is divided into an equal number of respiratory frames. The average time between respiratory frames is determined from the average respiration-rate and the number of desired respiration frames. When using this option, **window (%)** is neglected.

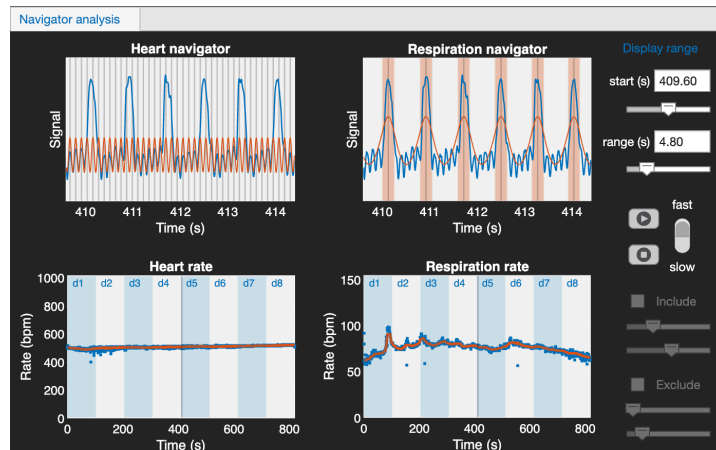
(4) Realtime

With the realtime option selected the data is sorted into 1 CINE movie per heart beat for fast dynamic imaging, for example during injection of a contrast agent. This option requires a short-as-possible TR and a suitable sparse k-space trajectory that results in sufficient k-lines per frame for compressed-sensing and/or parallel imaging image reconstruction. The number of dynamics will be automatically set to the total number of heartbeats in the acquisition and respiratory motion is disregarded. # frames under Frame sharing, in the Reco tab of panel 4, sets how many neighboring dynamics share k-lines, for better k-space filling and image quality at the expense of temporal averaging.

Press  to start sorting the data into k-space.

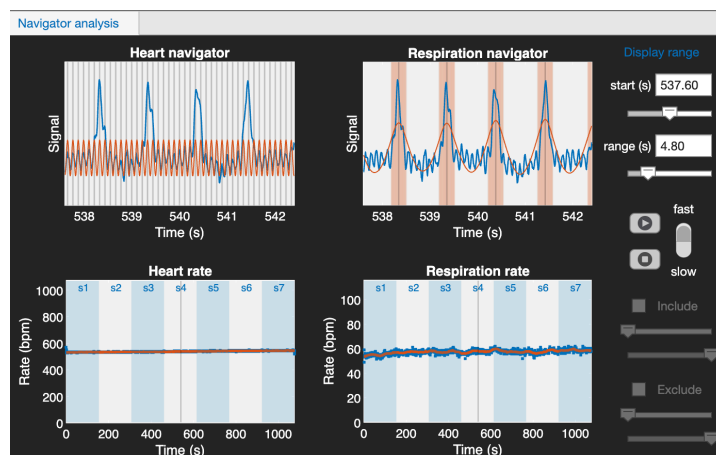
Dynamics

It is possible to reconstruct data in a dynamic time series, for example for dynamic-contrast enhanced acquisitions. Choose the desired number of dynamics, and continue with sorting k-space and reconstruction. There is no maximum number of dynamics, but remember that k-space will become very sparse with high number of dynamics and reconstruction might take very long. The dynamic time-frames will be indicated in the heart rate and respiratory rate plots in blue/white shaded areas as shown below.



Multi-slice 2D data

For multi-slice 2D acquisitions, the heart rate and respiration rate plots show when the subsequent slices were acquired with blue/white shaded areas. For multi-slice 2D data it is not possible to reconstruct different dynamics nor is it possible to include or exclude part of the scan.




3D data

Retrospective supports reconstructions of 3D CINE data.

Variable flip-angle data for T1-mapping

Cardiac quantitative T1 maps can be acquired using the variable flip angle method (see: Coolen et al. NMR Biomed 24, 154–162 (2011)). This option requires a 3D self-gated T1-weighted (RF and gradient spoiled) acquisition with variable flip angle (typically 5 angles, for example 2°, 5°, 10°, 15°, 20°). The varying flip angles should be acquired in a single acquisition in which the flip angles are evenly distributed over the k-space repetitions. After reconstruction data should be exported as a new MRD file for T1 mapping using the **T1mapp** application available on GitHub.

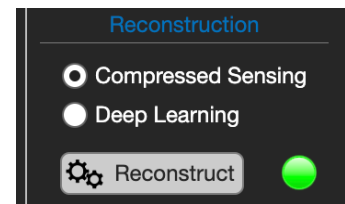




Step 6: Reconstruction of the data

Before reconstruction, select the appropriate reconstruction settings and regularization parameters in the **Reco** tab of panel 4.

Two reconstruction methods are available, selected in the Reconstruction panel of the dashboard:



(A) Compressed Sensing — the standard method, described below.

(B) Deep Learning — a trained neural network, described further down.

Press  to start the reconstruction.

A. Compressed Sensing reco parameters

Some parameters are not available without the BART toolbox.

(1) **WVxyz**

Wavelet regularization parameter in the spatial (x, y, and z) dimensions. For the mouse cardiac CINE, a low (0.001) or zero value resulted in the best movie quality.

(2) **TVxyz**

Total variation constraint in the spatial dimension. Higher values smooth the movie in the spatial domain. For mouse CINEs, this parameter does not seem to improve image quality. Recommended value therefore is zero.

(3) **LLRxyz**

Locally Low Rank regularization parameter; the block size is chosen from the data.

(4) **TVcine**

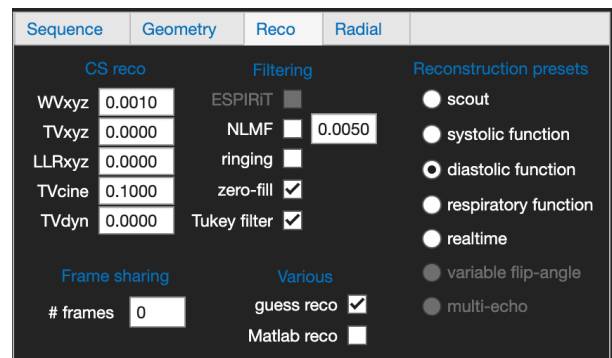
Total variation constraint in the CINE dimension. Higher values lead to improved image quality but increased temporal smoothing. For the mouse CINEs, a value of 0.1 seems a good compromise between image quality and temporal smoothing.

(5) **TVdyn**

Total variation constraint in the dynamics dimension.

(6) **ESPIRiT**

ESPIRiT reconstruction (Magn Reson Med. 2014 Mar 71(3): 990–1001.). This includes GRAPPA and SENSE reconstruction.





B. Deep learning reconstruction

Select Deep Learning in the Reconstruction panel. The deep learning reconstruction uses DRIM, a Dynamic Recurrent Inference Machine trained on rodent cardiac CINE data. Rather than solving an optimization problem as compressed sensing does, the network starts from the undersampled images and improves them in a fixed number of refinement steps, using the measured k-space at every step to keep the result consistent with the data that was actually acquired.

Requirements

The Python environment described in section 5.

macOS, or Windows with WSL.

2D Cartesian data, single or multi-slice, with at least 4 CINE frames; for 3D, radial, UTE and ZTE data the option is greyed out.

The app chooses the trained model automatically from the amount of undersampling, so there is nothing to set. The k-space filling percentage shown in the Sequence tab determines the choice: above about 67 % filling the model trained on lightly undersampled data is used, below that the model trained on more strongly undersampled data.

Points to be aware of

The network works on a 256 x 256 matrix, the matrix it was trained on; with zero-fill ticked in the Reco tab the images keep that size, otherwise they are resampled back to the acquired matrix. The reconstruction runs as a background process. No terminal window appears. The progress bar advances as the network iterates, and messages appear in the message box.

The network is applied once for every receive coil and every dynamic, so multi-coil or multi-dynamic data takes proportionally longer.

Memory use grows with the number of CINE frames, and the reconstruction time also with the number of slices; the matrix does not matter, since the network always works on 256 x 256. Before it starts, the app estimates the memory the reconstruction will need and compares it with what the graphics card, or the computer, offers; with CUDA, the card and shared memory together may be used, with a warning. When it will not fit, the reconstruction does not start: the message box says how much memory is needed, how much is available, and how many frames would fit. Fewer frames help; a smaller matrix does not.

Of the settings in the Reco tab of panel 4, the compressed sensing weights, ESPIRiT and Matlab reco do not apply to the deep learning reconstruction; NLMF and ringing are applied after it as well.

If the Python environment is not available, the deep learning option cannot be selected and the message box explains what is missing.



Sequence / Geometry

The Sequence and Geometry tabs contain information on the sequence parameters and the image geometry.

Sequence	Geometry	Reco	Radial
data file 7722_0.MRD			
freq Nx	187	FOV (mm)	35.00
phase Ny	192	TR (ms)	8.00
slices Nz	1	TE (ms)	2.00
coils Nc	1	averages	30.72
#repetitions	400	filling (%)	100
#flip angles	1	# k-lines	102400
flip angle(s)	20	scan time	00:13:39.200
		trajectory	user defined
		size	256
		time / dynamic	00:13:39.200

Sequence	Geometry	Reco	Radial
FOV (mm)		35.00	offset X (mm)
aspect ratio		1.00	offset Y (mm)
phase orientation		1	offset Z (mm)
slice thickness (mm)		1.00	angle X (deg)
slice gap (mm)		0.00	angle Y (deg)
			angle Z (deg)

Reco

The Reco tab contains additional reconstruction settings.

Sequence	Geometry	Reco	Radial
CS reco			
WVxyz	0.0010	ESPIRiT	☐
TVxyz	0.0000	NLMF	☐ 0.0050
LLRxyz	0.0000	ringing	☐
TVcine	0.1000	zero-fill	☒
TVdyn	0.0000	Tukey filter	☒
Reconstruction presets			
☐ scout ☐ systolic function ☒ diastolic function ☐ respiratory function ☐ realtime ☐ variable flip-angle ☐ multi-echo			
Frame sharing			
# frames	q		
Various			
guess reco ☒			
Matlab reco ☐			

Frame sharing: Binning of cardiac and/or respiratory motion states can be done using a Gaussian soft weighting of k-lines from neighboring frames (see: Küstner et al., NMR Biomed e4409 (2020). doi:10.1002/nbm.4409) to increase k-space filling and SNR. The #frames value determines the number of shared frames.

Reconstruction presets: Provides some pre-determined reconstruction settings for different applications.

NLMF: Applies a non-local means filter to the images with set weight.

ringing: Checkbox to switch an anti-ringing filter on or off.

zero-fill: Checkbox to switch zero-filling on or off.

Tukey filter: Checkbox to switch ellipsoid Tukey filtering of k-space on or off.

guess reco: Check the checkbox for the application to make an educated guess what type of reconstruction is appropriate. This is especially useful when doing automatic reconstruction.

Matlab reco: One can force a Matlab-based reconstruction by checking the checkbox.



Radial

The Radial tab contains additional settings for radial, UTE and ZTE reconstructions; ZTE scans have their own dead time (μs), shift and full echo.

Sequence	Geometry	Reco	Radial
Gradient delays calibration ☒ ring method ☐ Gx delay 0.00000 Gy delay 0.00000 Gz delay 0.00000 data offset 0 Radial trajectory ☐ half circle ☐ full circle ☒ user 1.00 center echo ☐ phase correction ☒ density correction ☒ ZTE dead time (μs) 14.0 shift 4 full echo ☐			

Tick the *calibration* checkbox for automatic adjustment of the gradient delays; it is ticked again at every load and Sort k-space. *Ring method* (2D radial, with Bart) uses Bart's *estdelay*, the method of Block and Uecker, not RING (Rosenzweig et al.), after which it is named; when it fails, it is unticked. When *ring method* is unchecked, a recursive image-based gradient delay optimization is performed, which is much slower.

The third delay field holds Gz for 3D UTE and the Sxy cross term for 2D radial data, and its label changes with the data. The delays are reset when a data set is loaded. For 3D UTE, with calibration ticked and Bart available, the data offset is found automatically too.


Three 2D radial trajectories are supported: half and full circle, with spokes spread evenly over 180° or 360° , and user, with the angle increment beside it, from the MRD file.

With *center echo* checked, the echoes are centered before reconstruction.

Phase correction performs a phase correction of the echoes before reconstruction.

Density correction performs a 2D or 3D density correction in k-space to account for non-uniform sampling. This setting only applies to reconstructions with *Bart*.

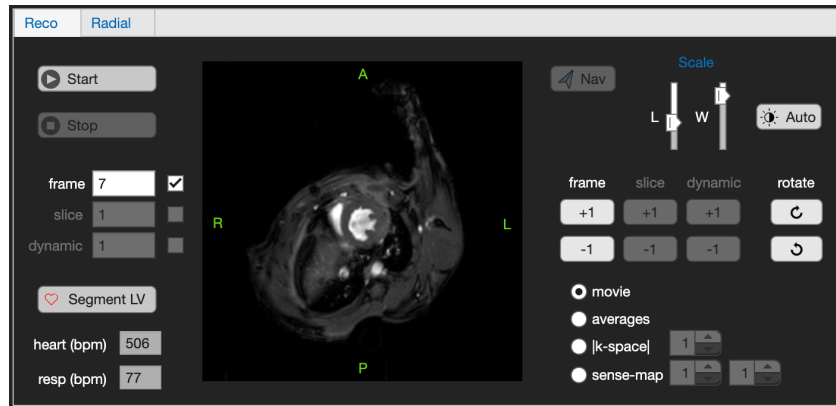
In the **Radial** tab in panel 5 one can view the 2D or 3D radial k-space filling and animate the 3D radial trajectory





Step 7: Viewing the movie

The resulting CINE MRI can be viewed in the Reco tab of panel 5. The radio buttons choose the movie, the number of averages per k-line and cardiac frame, the k-space, or the coil sensitivity maps (ESPIRiT).



Press  to start the movie.

Press  to stop the movie.

Check the corresponding boxes to loop the movie over frames, slices, and/or dynamics.

Image intensity can be adjusted with the level (**L**) and window (**W**) sliders.

Press  or  to rotate the movie clockwise or counter clockwise (this will not affect the Dicom and MRD export orientations).

Press  or  below **frame** to advance or rewind the movie 1 time frame.

Different slices can be viewed with  or  below **slice**.

Advance dynamics with  or  below **dynamic**.

The navigator button  toggles the visualization of the navigator slice (if visible in the current slice orientation)

Synchronizing multi-slice 2D CINEs from one animal

Since the cardiac synchronization is not based on ECG triggering, CINEs from different slices or acquired in different orientations may not be synchronized with respect to each other, i.e. the end-systole and end-diastole might have different frame numbers.

To correct for this the following procedure can be followed. First stop the movie with the  button. Then set the frame number to 1 in the dialog **frame** . Then, advance or rewind the movie with  or  to the desired starting phase of the cardiac cycle, for example end-diastole.

The cardiac synchronization will also be taken into account with DCM, GIF and MRD export.



Step 8: LV segmentation

Press Segment LV to outline the left ventricle in the reconstructed short-axis CINE images and calculate the ventricular volumes.

The segmentation uses a UNet3+ neural network, fine-tuned on rodent cardiac CINE images. It labels every pixel of every frame and every slice as one of three classes:

```
0    background
1    left ventricular myocardium
2    left ventricular blood pool
```

Requirements


Images must be reconstructed first.

The Python environment described in section 5.
macOS, or Windows with WSL.

Each dynamic is treated as a cardiac cycle of its own and is segmented separately. The network was trained on images with a fixed pixel size of 0.13672 mm on a 256 x 256 matrix, which corresponds to a field of view of 35 mm. The app resamples your images onto that grid automatically, so the matrix size and field of view of the acquisition do not matter in themselves. What does matter is that a field of view larger than 35 mm is cropped to the centre. If the heart is not near the centre of the image, part of it can be cut off. The message box reports the field of view against the 35 mm window and warns when it is exceeded.

Segmentation

Before starting the segmentation, make sure that all slices are cardiac phase aligned. See section "Synchronizing multi-slice 2D CINEs from one animal" on previous page.

To start the segmentation, press the  button.

Anatomical clean up


The raw network output is cleaned up with two rules based on anatomy:

- (1) The left ventricle is a single connected structure. Where a slice contains more than one separate object, only the one closest to the axis of the ventricle is kept.
- (2) The left ventricle lies on a consistent axis through the stack of slices. Where everything found in a slice lies too far from that axis, the ventricle has not been found and the slice is emptied. This is mostly what happens at the base and the apex.

The number of slices that were cleaned or emptied is reported in the message box and is written to the results file, so it can be checked afterwards.



Results



The results of the segmentation are shown in the Reco movie panel. Endocardium and epicardium will be shown as orange and blue contours, respectively.

The following are calculated per dynamic and shown in the message box:

EDV	end diastolic volume (ul)
ESV	end systolic volume (ul)
SV	stroke volume (ul)
EF	ejection fraction (%)
ED mass	left ventricular mass at end diastole (mg)

Where the slices do not touch, the slice separation rather than the slice thickness is used for the volume, so that the gaps are included. The myocardial mass is calculated with a density of 1.05 g/mL. The contours are drawn over the CINE movie: the epicardium is the outer border of the myocardium and the endocardium the border of the blood pool.

Settings: the entries of the settings file for the segmentation are listed in section 6.

Segmentation output files

The segmentation results are written by Export DCM (step 9), together with the exported images, so that the numbers and the contours travel with the data.

(1) Volumetry

LVsegmentation_<tag>.xlsx

An Excel file with two sheets.

Summary — one row per dynamic, with the columns:

```
DateTime, Tag, Dynamic, EDframe, ESframe, Frames, Slices,  
EDV_ul, ESV_ul, SV_ul, EF_pct, EDmass_mg,  
SliceSpacing_mm, ModelPixelSpacing_mm,  
CleanUp, DeblobbedSlices, ClearedSlices
```



PerFrame — the blood pool and myocardial volume of every frame:

Dynamic, Frame, BloodPool_ul, Myocardium_ul, IsED, IsES

This is the curve from which the ED and ES frames were chosen.

(2) Contours for Medis qMASS

`<PatientID>_STD<StudyID>_SER<n>_ACQ1_ITORIGINAL.con`

One file per dynamic, with `\_dyn<number>` added to the name when there is more than one. The file contains the endocardial and epicardial contours of every frame and every slice, and can be loaded onto the exported DICOM series.

(3) ROIs for Horos / OsiriX


`<PatientID>_STD<StudyID>_SER<n>_ACQ1_ITORIGINAL.rois_series`

The same contours in the format Horos uses for its own saved ROIs, written next to the exported DICOM files.


(4) Label maps as NiftI

`<name of the CINE file>_LVLabels.nii.gz`

Written next to the exported NiftI CINE, with the same header, so that it can be opened in ITK-SNAP as a segmentation of that image. The label values are the three classes listed at the start of this step.



ITK-SNAP



MEDIS



Step 9: Movie export

There are several ways to export the movies.



(1) Export DCM

Exports the data in Dicom and Nifti format for further processing in 3rd party software. The app searches for the Dicom information. If this information is not found, tags will be generated by the app itself. In the latter case the correct image position and orientation information are lost. The files go into a new numbered folder, `.../Data/<study>/DICOM/<series>R/<n>` and Nifti/`<tag>R/<n>`, or, without the scanner's Dicoms, into a folder that is asked for.

(2) Export GIF

Generates animated-gif movie or movies (3D, multi-slice) of the view on screen (movie, averages or k-space), with the LV contours when segmented. Image intensity (window, level) and image rotations are preserved. Gifs are saved in `.../Data/<study>/GIF/<tag>R/<n>`, or else in a folder you choose.

(3) Export MRD

The data will be stored as a new MRD file in the same directory as the input MRD file: `retro_<experiments>_<echoes>_<tag>.MRD`, with `retroData_..._f<filling>.mat`, which holds the k-space, images, averages and parameters. If the app finds an input rpr file, also a new rpr file will be generated. Additionally, when the app runs on the scanner computer, data will be reconstructed by the scanner software and new SUR and Dicom files will be placed in the data directory for viewing and planning of next scans in the MR Solutions preclinical software.

Export DCM also writes the heart and respiration navigators and the cardiac triggers as a BIDS physiological recording, `<tag>_physio.tsv.gz` with `<tag>_physio.json`, next to the NIfTI files. When the scanner's own Dicom images are not found, it asks where to save the export. It also writes the LV segmentation files (step 8).

Export MRD also writes the heart and respiration trigger times and rates as text files, next to the new MRD file. Every export also writes `recoparameters_<tag>.txt`.

Step 10: Exit

Press  **Exit** to shut down the Retrospective app. Closing the window does the same, also while the app is busy: it closes as soon as the running step has finished.



Credits

- BART toolbox and compressed sensing algorithms

<https://mrrecon.codeberg.page/>

Michael Lustig, David Donoho, John M. Pauly, Sparse MRI: The application of compressed sensing for rapid MR imaging, *Magn Reson Med* 58, 1182-1195 (2007).

- Automatic gradient delay estimation algorithm

Block, K. T. & Uecker, M., Simple method for adaptive gradient-delay compensation in radial MRI. *Proc. Intl. Soc. Mag. Reson. Med.* 19, 2816 (2011)

- Gibbs ringing suppression

Kellner, E, Dhital B., Kiselev VG and Reisert, M., Gibbs-ringing artifact removal based on local subvoxel-shifts. *Magn Reson Med*, 76(5), 1574-1581.

- NUFFT

<https://web.eecs.umich.edu/~fessler/>

- Phase unwrapping

M. A. Herráez, D. R. Burton, M. J. Lalor, and M. A. Gdeisat, "Fast two-dimensional phase-unwrapping algorithm based on sorting by reliability following a noncontinuous path", *Applied Optics*, Vol. 41, Issue 35, pp. 7437-7444 (2002),

H. Abdul-Rahman, M. Gdeisat, D. Burton, M. Lalor, "Fast three-dimensional phase-unwrapping algorithm based on sorting by reliability following a non-continuous path", *Proc. SPIE 5856, Optical Measurement Systems for Industrial Inspection IV*, 32 (2005).

- Frame sharing

Küstner et al., *NMR Biomed* e4409 (2020). doi:10.1002/nbm.4409

- Segmentation

Shah W, Stuckey DJ, Yao T, Wrobel M, Deniszczyc E, Feng Z, Anderson S, Campbell R, Muthurangu V, Steeden J. Open-source pre-clinical image segmentation: mouse cardiac magnetic resonance imaging datasets with a deep learning segmentation framework. *J Cardiovasc Magn Reson.* 2026 Summer;28(1):102706. doi: 10.1016/j.jocmr.2026.102706. Epub 2026 Feb 17. PMID: 41713653; PMCID: PMC13237537.